AI PRODUCT MANAGEMENT CASE STUDY · LIVE DEMO, LAUNCH CHECK PENDING · OCTOBER 2026

AI Personal Stylist

A lifelike video stylist that jots down what you want while you talk, then builds looks you refine by voice or touch

For anyone dressing for an occasion on their phone. A working web app, built in three weeks by an AI product manager directing Claude Code, with quality defined before it was measured. Sample products only.

97%
of notes jotted within 2 seconds of their phrase (target 95%, in-sample)

96%
of note checks right across 63 shopper sentences (target 90%, in-sample)

16 of 16
planted rule breaks caught by code before any look is shown

$1.15
of AI per 10-minute demo with styling changes: 5 cents over the $1.10 target

The problem

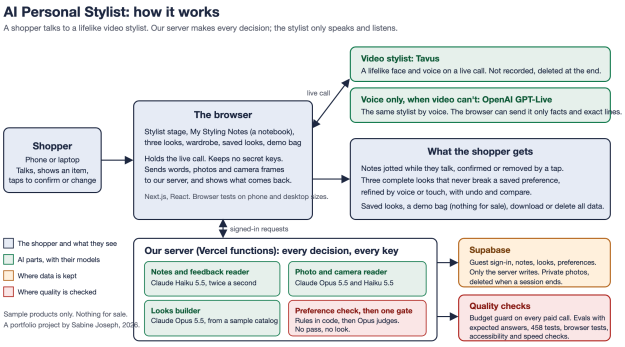
Online shopping gives you a search box and a thousand results that don't go together. A good stylist listens, remembers what you said, looks at what you own, and builds a whole outfit within your budget and rules. AI chat gets the listening half, then forgets your budget and shows you the heels you ruled out.

What I built

- **My Styling Notes:** the stylist jots free-form notes while you talk; a tap confirms or removes one, a correction by voice rewrites it. Nothing is a rule until confirmed.
- **Looks around what you own:** a photo or the live camera names your item; three complete looks are built around it, your own pieces never in the totals.
- **One gate:** every look passes a preference check (rules in code, then a model's judgment) before it reaches the screen.
- **Refinement and keeping:** keep, change, swap, undo, compare; save a look; a demo bag that cannot check out.
- **Degrade, never break; privacy and budget built in:** voice carries on when video can't; no recordings, photos deleted with the session, download or delete all; a daily allowance and monthly stops checked before every paid call.

Key decisions

- **A notebook, not a form:** notes in the stylist's own words, so the conversation feels heard; the price was a reader that must be fast and right mid-sentence.
- **The stylist only speaks and listens:** video and voice are swappable outside services; our server makes every decision and holds every key.
- **Accuracy first, then speed, then cost:** four models measured on the same sentences before any screen was built; the "best" one was too slow for live notes.
- **Specs before code:** product, design and tech specs and a phased plan, kept current; a phone check and a code review closed every phase.



The browser holds the call; our server reads notes and photos, builds looks and gates every one.

How I measured quality

Targets set on 17 to 21 October 2026, before any result, and never moved.

Measure	Target	Result
Notes within 2 s / right (63 sentences)	95% / 90%	97% / 96%
Rule breaks caught by code (37 looks)	All	16 of 16
Judged looks right	90%	21 of 21
Photo details right (24 catalog photos)	85%	93%
AI cost per demo	About $1.10	$1.15, not met
Phone speed, live welcome screen	90	98

What I learned

- Measuring before building changed the model choice, not just a setting.
- Cost became a product problem: one session's judgments hit the daily allowance in my own check, so the limit, honest limit lines and a free Undo followed.
- The automatic accessibility scan missed what measuring found: 200% text and 44 px targets caught five layout problems and one conflict in my own spec.

Limits: most answers were written by me, the builder, so the scores are in-sample; one blind set of 15 sentences is the exception (docs/evaluation.md). One tester, one iPhone; no stranger has used it yet. Catalog photos, not phone photos. Taste isn't measured, only rules and stated preferences.

In the repository: README, docs/evaluation.md, Project Memory (specs and plan), checks. Live at ai-personal-stylist-inky.vercel.app.

Skills: AI product management (specs from a PRD, a risk-first plan, metrics set before results, evaluation with its limits, model choice from data, cost and privacy by design), directing Claude Code as a non-developer.